
Instructions for paper submissions to AISTATS 2017

Anonymous Author 1
Unknown Institution 1

Anonymous Author 2
Unknown Institution 2

Anonymous Author 3
Unknown Institution 3

Abstract

1 Introduction

Reinforcement Learning is a paradigm where the learning algorithm/agent needs to learn via direct interaction with the environment. In most cases, the underlying environment can be modelled as a Markov Decision Process (MDP). At time any given time instant, the agent knows the state (described via the state variables) of the environment (i.e., the MDP) and has to choose an action from a given action set. The action selection mechanism of the agent is formally known as the policy, and each policy is associated with a value function, which is a mapping from the set of states to reals that specifies the value obtained by following a given policy starting from a given state. Typically, in order to improve a given policy, its value function needs to be computed first. Thus learning the value function of a given policy assumes centre stage.

Algorithms that learn the value function need to address two important issues namely lack of explicit model information (of the underlying MDP) and dimensionality of the state space. It is known that the value function of a policy can be computed by solving a $|S| \times |S|$ (where $|S|$ is the cardinality of the state space) linear system known as the Bellman equation (BE). However, the coefficients of the BE are described in terms of the model parameters, which are not known explicitly in an RL setting. Due to the *curse-of-dimensionality* $|S|$ grows exponentially in the number of state variables. As a result, when $|S|$ is very large, it is infeasible to store, retrieve and compute the value for each and every state.

Linear Function Approximation (LFA) is a common dimensionality reduction approach employed to

reduce the computational overhead when $|S|$ is large. In this approach, the value function is approximated by a linear combination of n ($\ll |S|$) basis functions. Once the basis is chosen, an approximate value function can be computed by solving a reduced ($n \times n$) linear system known as the projected Bellman equation (PBE). It is important to note that the coefficients of the PBE are dependent not only on the basis functions but also on the model parameters. Thus, RL algorithms need to compute the PBE solution using only the samples obtained via direct interaction.

The *temporal difference* (TD) family of learning algorithms are a very important sub-class of RL algorithms and compute the solution to the PBE. The term *temporal difference* stands for the error in the estimate of the value function obtained as a difference of terms involving the values at successive states. The first TD algorithm (i.e., TD(0)) was introduced in [1] and TD with LFA was analyzed in [2]. The TD family includes several variants such as TD such as TD(λ), SARSA, Q-learning, LSTD, iLSTD and Gradient-TD (GTD) [3]. Aspects such as sample efficiency, computational complexity and stability have driven the course of development of these various TD algorithms. The issue with TD(0) is that it makes use of only the current sample and is not sample efficient. This issue was alleviated by the LTSD by making use of all the data to estimate and solve an empirical PBE. Since LSTD involves matrix updates and inversions at each step its computational complexity is minimum $O(n^2)$. In contrast, since simple TD(0) involves only $O(n)$ computations per step, it continues to be of interest in applications in which cannot afford $O(n^2)$ required by LSTD.

The TD, Q-learning and SARSA algorithms are not stable in the most general setting, especially the *off-policy* case. Here, the policy (*target*) whose value function needs to be computed which is different from the policy (*behaviour*) that was used to collect the samples. The instability arises due to the fact that TD, SARSA and Q-learning are not true gradient algorithms. The Gradient TD (GTD) family of algorithms perform updates in the gradient direction and are sta-

ble. Even though GTD is stable in the off-policy case, it was pointed out in [1], that even the GTD is not a true stochastic gradient descent algorithm with respect to its original objective [1]. Authors in [2] propose newer variants of the GTD and show that they are true gradient algorithms, however, with respect to a different objective function obtained via a primal-dual saddle point formulation.

Another related and important aspect of concern is the speed of convergence. The idea of speeding TD learning was used in speedy Q-learning (SQL) [3]. [3] uses adaptive step sizes to achieve faster learning rates. Recent RL algorithms such as GTD-MP claim to make use of ideas of accelerated gradient descent to achieve increase speed of convergence.

Motivation Efforts in the past have been to adopt a different horse for different course approach, i.e, methods have been developed to address one issue (computational/speed/stability) at a time. This approach has led to ad-hoc derivation of the updates, speeding techniques and step-size rules, thus does not provide a complete holistic picture. The focus of this paper is to study these TD algorithms in a more principled and unified manner using the theory of stochastic approximation (SA) algorithms. The term stochastic approximation signifies the fact that the noisy quantities used in the algorithms are stochastically approximate, and are equal to the quantities only in the expectation. Since the TD algorithms are also based on such stochastic updates, we make use of recent finite time results in SA [4] to argue our case. The specific contributions in this paper are listed as below

- We use of the standard result in SA theory to identify ordinary differential equations related to each TD algorithm. Thus each of the different TD algorithms can be viewed as a discretization of its corresponding ODE. While all the ODEs converge to the PBE solution, their dynamics is dictated by different design matrices. Thus choosing design matrices that lead to stable ODEs is key.
- The finite time performance of any TD algorithm has two terms namely the bias term (due to the initial conditions) and the variance terms (due to the noisy samples). Further, we then argue on the lines of [4] to show that using conventional step size rules that diminish in $O(1/n)$ are sensitive to the spectral properties of the design matrix.
- We then show that constant step size with Rupert-Polyak averaging of iterates yields $O(1/n^2)$ for forgetting initial conditions and $O(1/n)$ rate for variance term.
- The expression for finite time performance (initial

condition and bias) involves constant terms which depend on the spectral properties of the design matrix. For instance, the ‘slowness of GTD can be attributed to its design matrix.

- Our SA based approach is not only useful for analysis but also for synthesis of RL algorithms. For instance, we show that the GTP-MP algorithm can be derived as a predictor-corrector discretization of the ODE corresponding to the GTD algorithm. We add more weight to this argument by deriving a new stable TD method by modifying the predictor-corrector algorithm.

We also present numerical examples along side the analysis.

2 Background

We now present a brief overview of Markov Decision Processes (MDPs) which is a framework to describe the RL setting, i.e., the agent’s interaction with the environment. An MDP is a five tuple and is denoted by $M = \{S, A, P, R, \gamma\}$, where S is the state space, A is the action space, $P = (p_a(s, s'), a \in A, s, s' \in S)$ is the probability transition kernel that specifies the probability of transitioning from state s to s' when the agent taking an action a , and $R(s, a): S \times A \rightarrow \mathbb{R}$ is the reward for taking an action a in state s and $0 < \gamma < 1$ is a given discount factor. A stationary deterministic policy (or simply a policy) is denoted by $\pi = (\pi(s, \cdot), s \in S)$, where $\pi(s, \cdot)$ is a probability distribution over the set of actions. The infinite horizon discounted reward of a policy π is given by

$$J_\pi = \mathbf{E}[\sum_{t=0}^{\infty} \gamma^t R(s_t, a_t) | s_0 = s, \pi], \quad (1)$$

and is the discounted sum of rewards starting from state s and taking actions according to policy π .

It is of interest to compute the optimal policy π^* which yields the maximum reward starting from any given state. A key step in computing the optimal policy is *policy evaluation*, i.e., computing J_π of a given policy π .

2.1 Bellman Equation

The value function J_π can be computed by solving the Bellman equation (BE) given below:

$$J_\pi = T_\pi J_\pi, \quad (2)$$

where in (4) $T_\pi: \mathbb{R}^{|S|} \rightarrow \mathbb{R}^{|S|}$ is the Bellman operator. The Bellman operators is defined in terms of the model

parameters of the MDP as follows

$$(T_\pi J)(s) = \sum_{a \in A} \pi(s, a) (R(s, a) + \gamma \sum_{s' \in S} p_a(s, s') J(s')), \quad (3)$$

where $J_\pi = (J_\pi(s), \forall s \in S)$. The BE can thus be re-written as

$$J_\pi = R_\pi + \gamma P_\pi J_\pi, \quad (4)$$

where $R_\pi(s) = \sum_{a \in A} \pi(s, a) R(s, a)$ and $p_\pi(s, s') = \sum_{a \in A} \pi(s, a) p_a(s, s')$. Note that the BE is a linear system of equations in $|S|$ variables.

2.2 Function Approximation

The term *curse-of-dimensionality* denotes the fact that $|S|$ scales exponentially in the number of state variables. Due to the curse, MDPs in practice tend to have large number of states. Solving the linear system in (4) involves matrix inversion and requires a minimum of $O(|S|^2)$ computations, and hence it is impractical to solve the BE when $|S|$ is large. A practical way to address this issue is to approximate the value function. A widely adopted approach is function approximation, wherein the value function is approximated by a parameterized family of functions $J_\pi \approx J_\theta = \Phi(\theta^*)$, where $\{\Phi(\theta), \theta \in \mathbb{R}^n\}$ is the parameterized family and $\theta^* \in \mathbb{R}^n$ is the optimal parameter. Note that the parameter θ belongs to \mathbb{R}^n and dimensionality reduction can be achieved by choosing $n \ll |S|$. Linear parameterization is the most common, wherein, $J_\pi \approx J_\theta \Phi \theta^*$, with Φ as a $|S| \times n$ feature matrix.

2.3 Projected Bellman Equation

Once Φ is fixed, we have to come up with a procedure to compute out θ^* . An immediate idea is to take a cue from linear regression and look at the following quadratic loss function that quantifies the approximation error.

$$L(\theta) = \sum_{s \in S} (J_\theta - T J_\theta)^2. \quad (5)$$

The above penalty function is called *mean-square-Bellman-error* (MSBE) and a way to compute θ^* would be to minimize the MSBE. A caveat however is that $T J_\theta$ does not usually belong to the subspace belonging to spanned by the columns of Φ . As a result, $T J_\theta$ cannot be expressed as $\Phi \theta$ for any $\theta \in \mathbb{R}^n$. Due to this reason, none of the TD algorithms consider the MSBE as the error criterion.

A way to side step the issue of non-representability of $T J_\theta$, we can instead consider a projected Bellman

equation (PBE) given by

$$\begin{aligned} J_{\theta^*} &= \Pi T_\pi J_{\theta^*}, \\ \Phi \theta^* &= \Pi T_\pi \Phi \theta^*, \end{aligned} \quad (6)$$

where $\Pi = \Phi(\Phi^\top D_\pi \Phi)^{-1} \Phi^\top$ is the least squares projection operator and D_π is a diagonal matrix with the diagonal entry $D_\pi(s, s)$ as the stationary distribution of state s under policy π .

3 Stochastic Approximation

A typical single timescale stochastic recursion is given by

$$x_{t+1} = x_t + \alpha_t (f(x_t) + M_{t+1}), \quad (7)$$

where $x_t \in \mathbb{R}^n, \forall t \geq 0$, $f: \mathbb{R}^n \rightarrow \mathbb{R}^n$ is a Lipschitz continuous function and $M_{t+1}, t \geq 0$ are martingale difference noise terms. Under reasonable conditions, we can analyze the iterates in (7) by studying its average behaviour. To this end, in addition to Lipschitz continuity of f , we need following assumptions.

Assumption 1.

1. $\alpha_t \rightarrow 0, t \rightarrow \infty, \sum_{t \geq 0} \alpha_t = \infty$ and $\sum_{t \geq 0} \alpha_t^2 < \infty$.
2. $\{M_t\}$ is a martingale difference sequence with respect to an increasing family of σ -fields $\mathcal{F}_t \doteq \sigma(x_0, M_1, \dots, M_t), t \geq 0$, such that $\mathbf{E}[M_{t+1} | \mathcal{F}_t] = 0$ a.s., $t \geq 0$. Further, $\{M_t\}$ are square-integrable i.e., $\mathbf{E}[M_{t+1}^2 | \mathcal{F}_t] \leq K(1 + \|x_t\|^2)$, a.s., $\forall t \geq 0$.
3. $\sup_t \|x_t\| < \infty$
4. The ordinary differential equation (ODE) $\dot{x}(t) = f(x(t))$, has a globally unique asymptotically stable equilibrium $x^* \in \mathbb{R}^n$.

Equipped with the above assumptions we state an useful result from []

Theorem 1. Under the above assumptions we have $x_t \rightarrow x^*$ as $t \rightarrow \infty$.

The algorithms of interest in this paper can be written as linear updates that assume the following form:

$$x_{t+1} = x_t + \alpha_t (g - H x_t), \quad (8)$$

where $x_t \in \mathbb{R}^n$ are the iterates, $H \in \mathbb{R}^{n \times n}$ the design matrix and $g \in \mathbb{R}^n$ is the design vector. In what follows, we use the notation (g, H) to specify the design of the linear recursion in (8). The following corollary is a direct consequence of Assumption 1 4 and Theorem 1

Corollary 1. The iterates $x_t \in \mathbb{R}^n$ in the linear stochastic update rule in (8) converge $x^* = H^{-1}g$ as $t \rightarrow \infty$ iff all the Eigen values of H have positive real parts.

An important point to note is that the linear update rule in (8) can be viewed as an Euler discretization of the ODE $\dot{x}(t) = (g - Hx(t))$.

4 Temporal Difference Learning Algorithms

The *temporal difference* (TD) class of learning algorithms solve the projected Bellman equation. The samples of the TD family namely TD(0) is given by,

$$\theta_{t+1} = \theta_t + \alpha_t \underbrace{(R(s_t, a_t) + \gamma \phi(s_{t+1})\theta_t - \phi(s_t)\theta_t)}_{\delta_t} \phi(s_t)^\top, \quad (9)$$

where $\phi(s_t)^\top \in \mathbb{R}^n$ is the feature of state s_t and is the s^{th} row of the feature matrix Φ . In (10), δ_t is the temporal difference term which is an estimate of the error in the value function specified by θ_t . The TD(0) update can also be re-written as

$$\theta_{t+1} = \theta_t + \alpha_t (b_\pi - A_\pi) \theta_t + \alpha_t M_{t+1} \quad (10)$$

where $b_\pi = \Phi^\top D_\pi R_\pi$ and $A_\pi = \Phi^\top D_\pi (I - \gamma P_\pi) \Phi$. In order to arrive at (10) we have made use of the fact that $s_t \sim D_\pi, a_t \sim \pi(s_t, \cdot)$, $\mathbf{E}[\delta_t | s_t] = (R_\pi(s_t) + \gamma(P_\pi J_\theta)(s_t) - J_\theta(s_t))$ and letting $M_{t+1} = \delta_t \phi(s_t)^\top - \mathbf{E}[\delta_t \phi(s_t)^\top]$.

Note that the design of TD(0) is given by (b_π, A_π) .

4.1 Off-policy divergence of TD(0)

The TD(0) algorithm can exhibit divergent behaviour in the *off-policy* setting. In the *off-policy* setting the recursion in (10) has to be modified as follows

$$\theta_{t+1} = \theta_t + \alpha_t (\rho_t(s_t, a_t) R(s_t, a_t) + \gamma \phi(s_{t+1})\theta_t - \phi(s_t)\theta_t) \phi(s_t)^\top, \quad (11)$$

where $\rho_t = \frac{\pi(s_t, a_t)}{\mu(s_t, a_t)}$ is the importance sampling ratio between the behaviour policy μ and target policy π . The recursion in (11) can be re-written as

$$\theta_{t+1} = \theta_t + \alpha_t (b_\pi - A_\pi^\mu) \theta_t + \alpha_t M_{t+1}^\mu, \quad (12)$$

where $A_\pi^\mu = \Phi^\top D_\mu (I - \gamma P_\pi) \Phi$, where μ is the behaviour policy used to sample the data. Thus the design of TD(0) can in the off-policy case is (b_π, A_π^μ) . The undoing of TD(0) in the off-policy is due to the fact that all the eigen values of A_π^μ cannot be guaranteed to have positive real parts. This condition can be demonstrated via the following simple example from [1].

Example 1. Consider a two state MDP with $P_\pi = \begin{bmatrix} 0 & 1 \\ 0 & 1 \end{bmatrix}$, $\Phi = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$, $D_\mu = \begin{bmatrix} 0.5 & 0 \\ 0 & 0.5 \end{bmatrix}$. Then the design matrix $A_\pi^\mu = \Phi^\top (I - \gamma P_\pi) \Phi = -0.2$ is negative definite and hence from Corollary 1 it follows that TD(0) is divergent for this example.

4.2 Gradient Temporal Difference Learning

The instability of TD(0) is due to the fact that it is not a true gradient descent algorithm. The first gradient-TD (GTD) algorithm was proposed in [2] and is based on minimizing the *norm of the expected TD update* (NEU) given by

$$NEU(\theta) = \mathbf{E}[\rho \phi^\top \delta(\theta)] \mathbf{E}[\rho \phi^\top \delta(\theta)] \quad (13)$$

The GTD scheme based is on the gradient of the above expression which is given by $-\frac{1}{2} \nabla NEU(\theta) = \mathbf{E}[\rho(\phi - \gamma \phi)^\top \phi] \mathbf{E}[\rho \phi^\top \delta(\theta)]$. Since the gradient is a product of two expectation we cannot use a sample product (due to the presence of correlations). The GTD addresses this issue by estimating $\mathbf{E}[\rho \delta \phi^\top]$ in a separate recursion. The GTD updates can be given by

$$\begin{aligned} y_{t+1} &= y_t + \alpha_t (\rho_t \phi^\top \delta_t - y_t) \\ \theta_{t+1} &= \theta_t + \alpha_t \rho_t (\phi(s_t) - \gamma \phi(s_{t+1}))^\top \phi(s_t) y_t \end{aligned} \quad (14)$$

The design of GTD is given by (g_1, H_1) where $g_1 = \begin{bmatrix} b_\pi \\ 0 \end{bmatrix}$, $H_1 = \begin{bmatrix} -I & -\Phi^\top D_\mu (\Phi - \gamma P_\pi \Phi) \\ (\Phi - \gamma P_\pi \Phi)^\top D_\mu \Phi & 0 \end{bmatrix}$.

Instead of NEU, the *mean-square projected Bellman Error* (MSPBE) can also be minimized. The MSPBE is defined as

$$MSPBE(\theta) = \| J_\theta - \Pi T_\pi J_\theta \|_D^2 \quad (15)$$

The GTD2 algorithm was proposed in [3] based on minimizing (15). The GTD2 updates are given by

$$\begin{aligned} y_{t+1} &= y_t + \beta_t \phi(s_t)^\top (\rho_t \delta_t - \phi(s_t) y_t) \\ \theta_{t+1} &= \theta_t + \alpha_t \rho_t (\phi(s_t) - \gamma \phi(s_{t+1}))^\top \phi(s_t) y_t \end{aligned} \quad (16)$$

The design of GTD2 is give by $g_2 = (b_\pi, 0)$, $H_2 = \begin{bmatrix} -\Phi^\top D_\mu \Phi & -\Phi^\top D_\mu (\Phi - \gamma P_\pi \Phi) \\ (\Phi - \gamma P_\pi \Phi)^\top D_\mu \Phi & 0 \end{bmatrix}$

4.3 Speedy Learning

Speedy Q-learning was introduced in [4] and uses a ‘speed-up term to achieve better convergence rates. The SQL updates for the full state representation is given by

$$\begin{aligned} J_{t+1}(s) &= J_t(s) + \alpha_t (R(s_t, a_t) + \gamma J_{t-1}(s_{t+1}) - J_t(s_t)) + \\ & (1 - \alpha_t) ((R(s_t, a_t) + \gamma J_t(s_t)) - (R(s_t, a_t) + \gamma J_{t-1}(s_t, a_t))) \end{aligned} \quad (17)$$

The design corresponding to SQL is $((I - \gamma P_\pi)^{-1} R_\pi, I)$.

5 Finite Time Analysis

We now study the finite time performance of the linear SA update in (8)

$$\begin{aligned} x_{t+1} - x^* &= x_t + \alpha_t(g - Hx_t) + \alpha_t(M_{t+1}) - H^{-1}g \\ &= (I - \alpha_t H)(x_t - x^*) + \alpha_t M_{t+1} \\ &= \underbrace{\prod_{i=0}^t (I - \alpha_i H)(x_0 - x^*)}_{\text{Bias}} \\ &\quad + \underbrace{\sum_{i=0}^t \prod_{i+1 \leq k < t} (I - \alpha_k H) \alpha_i M_{i+1}}_{\text{Noise}} \end{aligned}$$

5.1 Conventional step-size rule

We now show sensitivity of the rate of convergence to the spectral properties of the design matrix in the case of conventional step size namely $\alpha_t = \frac{C}{t}$. To this end, we assume that all the Eigen values $\mu_i, i = 0, \dots, n-1$ of H have positive real parts and that μ_0 is an Eigen value such that $|(1 - \alpha_t \mu_0)| \leq |(1 - \alpha_t \mu_i)|, \forall t \geq 0, i = 1, \dots, n-1$. We now write down the expressions for the bias and the variance terms separately.

$$\begin{aligned} &\prod_{i=0}^t (I - \alpha_i H)(x_0 - x^*) + \sum_{i=0}^t \prod_{i+1 \leq k < t} (I - \alpha_k H) \alpha_i M_{i+1} \\ &\leq \exp^{-C\mu \sum_{i=0}^t \alpha_i} |x_0 - x^*| + \exp^{-2C\mu t} \sum_{k=1}^t \exp(2C\mu \log k) \alpha_k^2 |M_{k+1}|^2 \\ &\leq \exp^{-C\mu \log t} |x_0 - x^*| + \exp^{-2C\mu t} \frac{t^{2C\mu-2+1}}{2C\mu-2+1} C^2 \sigma^2 \\ &\leq \frac{1}{t^{C\mu}} |x_0 - x^*| + \frac{1}{t} \frac{C^2}{2C\mu-1} \sigma^2 \end{aligned} \quad (18)$$

Notice that while the noise term reduces $O(1/t)$, the bias term vanishes only at a rate $O(1/t^{C\mu})$ and shows sensitivity to the spectral property of the design matrix.

5.2 Rupert-Polyak Averaging with constant step size

Let $\alpha_t = \alpha$ in (??) and then consider $\bar{x}_t = \frac{1}{t+1} \sum_{k=0}^t x_k$, we then have

$$\bar{x}_t - x^* = \frac{1}{t+1} \left(\sum_{i=0}^t (I - \alpha H)^{t-i} (x_0 - x^*) + \alpha \sum_{i=0}^t \sum_{k=0}^{t-i} (I - \alpha H)^{t-i-k} M_{i+1} \right) \quad (25)$$

Let $\bar{\lambda}$ be the eigen value of $(I - \alpha H)$ with maximum modulus, then we have

$$\mathbf{E}[\|\bar{x}_t - x^*\|^2] \leq \frac{1}{(t+1)^2} \frac{1}{(1 - \bar{\lambda})^2} (\|x_0 - x^*\|^2 + \alpha^2 \sum_{i=0}^t \mathbf{E} \|M_{i+1}\|^2) \quad (19)$$

6 Numerical Illustration

In this section, we compute the H matrix for the various TD algorithms discussed so far and make qualitative and quantitative comments. In order to keep the discussion simple and intuitive, we compute explicit values for the most trivial MDP namely single state MDP.

7 A new stable TD algorithm

We now modify the predictor corrector algorithm as follows

$$x_t^m = x_t + \beta(g - Hx_t) \quad (20)$$

$$x_{t+1} = x_t + \alpha_t(g - Hx_t^m) \quad (21)$$

The above recursion can then be written as

$$x_{t+1} = (I - \alpha_t(H + \beta H^2))x_t + \alpha_t(g - \beta Hg) \quad (22)$$

Notice that the matrix $H + \beta H^2$ can always be made have all real values positive.

Example 2.

8 Synthesis of New algorithms

8.1 Stories of Failures

8.2 Stories of Success

Consider the ODE given by

$$\dot{x}(t) = (g - Hx(t)) \quad (23)$$

with initial condition $x(0) = x_0$. The Euler discretization of the above ODE leads to the following update rule

$$x_{t+1} = x_t + \alpha_t(g - Hx_t). \quad (24)$$

The *Predictor-Corrector* discretization of the ODE is the following update rule

$$x_t^m = x_t + \alpha_t(g - Hx_t) \quad (25)$$

$$x_{t+1} = x_t + \alpha_t(g - Hx_t^m) \quad (26)$$

The idea behind the PC method is to first take a step in the gradient direction to produce a new point x_t^m , and then obtain the estimate of the gradient at x_t^m to be used to update x_t . It is clear that the GTD-MP is using the PC type update. Notice that the PC updates can be unfurled and written as the following single recursion

$$x_{t+1} = x_t + \alpha_t(g - H(x_t + \alpha_t(g - Hx_t))) \quad (27)$$

$$= x_t + \alpha_t(g + \alpha_t g - H + \alpha_t H^2)x_t \quad (28)$$

$$= (I - \alpha_t H + \alpha_t^2 H^2)x_t + \alpha_t(g - \alpha_t Hg) \quad (29)$$

Notice that the PC method is slower than the Euler discretization in (24).